

```
In [304... from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.decomposition import PCA
import numpy as np
from sklearn.metrics import precision_score, recall_score, f1_score
from RevGEN_MLP import RevGEN_MLP
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import gaussian_kde
from matplotlib.colors import ListedColormap
from sklearn.ensemble import IsolationForest, RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_score
import torch
import torch.nn as nn
import torch.optim as optim
from sklearn.metrics import accuracy_score, classification_report
```

## RevGEN-MLP

### Loading and Preparing Data for Training

```
In [305... df = pd.read_csv("wine+quality\\winequality-white.csv", sep=";") # Insert path to dataset here
df['good'] = (df['quality'] >= 7).astype(int)
X = df.drop(['quality', 'good'], axis=1)
y = df['good']

column_names = X.columns

X = X.values
y = y.values
```

```
In [306... # Split into train/test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
# One-hot encode labels
encoder = OneHotEncoder(sparse_output=False)
y_train_encoded = encoder.fit_transform(y_train.reshape(-1, 1))
y_test_encoded = encoder.transform(y_test.reshape(-1, 1))

# Scale features
ANN_scaler = StandardScaler()
X_train_scaled = ANN_scaler.fit_transform(X_train)
X_test_scaled = ANN_scaler.transform(X_test)
```

### Training RevGEN-MLP

```
In [307... num_layer = 3
num_epochs = 101
```

```
In [308... # Set seed for reproducibility
np.random.seed(42)

# Initialize model
model = RevGEN_MLP(
    n_layers=num_layer,
    x=X_train_scaled[0].reshape(-1, 1),
    y_actual=y_train_encoded[0].reshape(-1, 1),
    epochs=num_epochs,
    loss_function="cross_entropy"
)

# Training Loop
for epoch in range(num_epochs):
    loss_epoch = 0
    correct_train = 0

    # Shuffle training indices
    indices = np.random.permutation(X_train_scaled.shape[0])

    for i in indices:
        x_sample = X_train_scaled[i].reshape(-1, 1)
        y_sample = y_train_encoded[i].reshape(-1, 1)

        # Train and accumulate loss
```

```

model.train(input=x_sample, target=y_sample)
loss_epoch += model.loss_fn(input=x_sample, target=y_sample)

# Predict and count correct predictions
pred = model.forward(x_sample)
if np.argmax(pred[:2]) == np.argmax(y_sample):
    correct_train += 1

# Compute average training metrics
avg_train_loss = loss_epoch / X_train_scaled.shape[0]
train_accuracy = correct_train / X_train_scaled.shape[0]

# Evaluate on test set every 10 epochs
if epoch % 10 == 0 or epoch == num_epochs - 1:
    correct_test = 0
    test_loss_epoch = 0

    for i in range(X_test_scaled.shape[0]):
        x_sample = X_test_scaled[i].reshape(-1, 1)
        y_sample = y_test_encoded[i].reshape(-1, 1)

        pred = model.forward(x_sample)
        if np.argmax(pred[:2]) == np.argmax(y_sample):
            correct_test += 1

        test_loss_epoch += model.loss_fn(input=x_sample, target=y_sample)

    avg_test_loss = test_loss_epoch / X_test_scaled.shape[0]
    test_accuracy = correct_test / X_test_scaled.shape[0]

    print(f"Epoch {epoch:3d} | "
          f"Train Loss: {avg_train_loss:.4f} | Train Acc: {train_accuracy * 100:.2f}% | "
          f"Test Loss: {avg_test_loss:.4f} | Test Acc: {test_accuracy * 100:.2f}%")

```

Epoch	0	Train Loss: 0.5016	Train Acc: 77.44%	Test Loss: 0.4701	Test Acc: 78.16%
Epoch	10	Train Loss: 0.3573	Train Acc: 81.60%	Test Loss: 0.4194	Test Acc: 78.16%
Epoch	20	Train Loss: 0.3402	Train Acc: 83.51%	Test Loss: 0.4025	Test Acc: 80.82%
Epoch	30	Train Loss: 0.3296	Train Acc: 83.97%	Test Loss: 0.3938	Test Acc: 82.24%
Epoch	40	Train Loss: 0.3293	Train Acc: 83.92%	Test Loss: 0.3862	Test Acc: 81.94%
Epoch	50	Train Loss: 0.3202	Train Acc: 84.64%	Test Loss: 0.3898	Test Acc: 81.63%
Epoch	60	Train Loss: 0.3135	Train Acc: 85.17%	Test Loss: 0.3814	Test Acc: 82.35%
Epoch	70	Train Loss: 0.3057	Train Acc: 85.40%	Test Loss: 0.3785	Test Acc: 82.35%
Epoch	80	Train Loss: 0.2987	Train Acc: 85.71%	Test Loss: 0.3795	Test Acc: 82.55%
Epoch	90	Train Loss: 0.2936	Train Acc: 86.29%	Test Loss: 0.3792	Test Acc: 82.55%
Epoch	100	Train Loss: 0.2876	Train Acc: 86.65%	Test Loss: 0.3653	Test Acc: 83.37%

In [309...

```

correct = 0
y_preds = []
y_true = []

for i in range(X_test_scaled.shape[0]):
    x_sample = X_test_scaled[i].reshape(-1, 1)
    y_sample = y_test_encoded[i].reshape(-1, 1)

    pred = model.forward(x_sample)
    pred_class = np.argmax(pred[0:2])
    true_class = np.argmax(y_sample)

    y_preds.append(pred_class)
    y_true.append(true_class)

    if pred_class == true_class:
        correct += 1

x_sample = X_test_scaled[0].reshape(-1, 1)
pred = model.forward(x_sample)

test_accuracy = correct / X_test_scaled.shape[0]
precision = precision_score(y_true, y_preds, average='macro', zero_division=0)
recall = recall_score(y_true, y_preds, average='macro', zero_division=0)
f1 = f1_score(y_true, y_preds, average='macro', zero_division=0)

print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
print(f"Precision: {precision * 100:.2f}%")
print(f"Recall: {recall * 100:.2f}%")

```

Test Accuracy: 83.37%  
Precision: 75.82%  
Recall: 72.14%

## Invertibility Check

Small reconstruction errors close to zero are expected due to floating-point precision limits

```
In [310... x_sample = X_test_scaled[0].reshape(-1, 1)

# Unscale
x_sample_unscaled = ANN_scaler.inverse_transform(x_sample.reshape(1, -1))
print("Original Sample:", x_sample_unscaled)

# Forward pass
pred = model.forward(x_sample)
print("\nOutput (Classes):", pred[0:2].ravel())
print("Output (Latent Variable):", pred[2:].ravel())
# Reconstruct
reconstructed_sample = model.reverse(pred)

# Unscale reconstruction
reconstructed_sample_unscaled = ANN_scaler.inverse_transform(
    reconstructed_sample.reshape(1, -1)
)
print("\nReconstructed Sample:", reconstructed_sample_unscaled.ravel())

mse_scaled = np.mean((x_sample - reconstructed_sample)**2)
mse_unscaled = np.mean((x_sample_unscaled - reconstructed_sample_unscaled)**2)
print("\nMSE Error (Scaled Data): ", mse_scaled)
print("MSE Error (Unscaled Data): ", mse_unscaled)
```

Original Sample: [[6.0000e+00 1.7000e-01 3.6000e-01 1.7000e+00 4.2000e-02 1.4000e+01  
6.1000e+01 9.9144e-01 3.2200e+00 5.4000e-01 1.0800e+01]]

Output (Classes): [0.98050259 0.01949741]

Output (Latent Variable): [ 3.92619796e+00 -4.75037097e-02 -3.29723407e-02 -1.94034664e-02  
-3.49284453e-02 -1.56047654e-03 -4.75244660e-02 9.84343548e-01  
-6.25475829e-02]

Reconstructed Sample: [6.00000000e+00 1.70000000e-01 3.60000000e-01 1.70000002e+00  
4.20000001e-02 1.39999999e+01 6.10000002e+01 9.91440000e-01  
3.22000000e+00 5.39999999e-01 1.08000000e+01]

MSE Error (Scaled Data): 2.2025489842542087e-17

MSE Error (Unscaled Data): 3.1172022567704e-15

## Generation

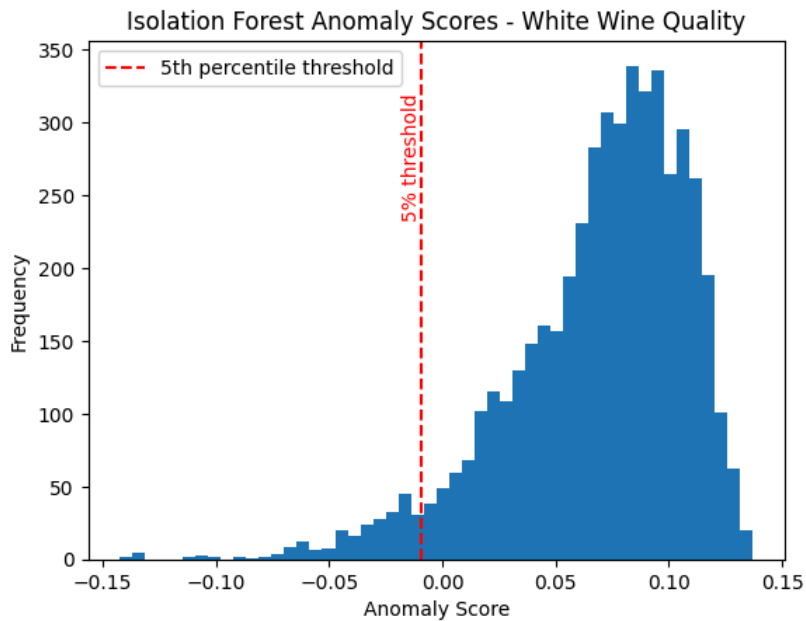
```
In [311... detector = IsolationForest(random_state=42).fit(X)

scores = detector.decision_function(X)

# Plot histogram
plt.hist(scores, bins=50)
threshold = np.percentile(scores, 5)
plt.axvline(threshold, color='red', linestyle='--', label='5th percentile threshold')

# Add label
plt.text(threshold, plt.ylim()[1]*0.9, '% threshold', color='red', rotation=90, va='top', ha='right')

plt.title("Isolation Forest Anomaly Scores - White Wine Quality")
plt.xlabel("Anomaly Score")
plt.ylabel("Frequency")
plt.legend()
plt.show()
```



```
In [312... def generate_data_with_epsilon(model, epsilon):
    pred_classes = []
    inputs = []
    original_inputs = []
    original_classes = []
    original_prob = []
    pred_probability = []

    for test, truth in zip(X_test_scaled, y_test_encoded):
        x_sample = test.reshape(-1, 1)
        pred = model.forward(x_sample)

        class_probs = pred[0:2].ravel()
        max_prob = np.max(class_probs)
        if np.argmax(pred[:2]) == np.argmax(truth):
            if max_prob >= 0.8:
                original_prob.append(max_prob)
                pred_class = np.argmax(class_probs)
                original_inputs.append(ANN_scaler.inverse_transform(x_sample.reshape(1, -1))[0])
                original_classes.append(pred_class)
                vector = pred.copy()
                vector[2:] = vector[2:] + epsilon

                new_input = model.reverse(input=vector)
                new_pred = model.forward(new_input)
                max_prob = np.max(new_pred[0:2])
                pred_probability.append(max_prob)
                new_pred_class = np.argmax(new_pred[0:2])

                new_input_unscaled = ANN_scaler.inverse_transform(new_input.reshape(1, -1))[0]

                pred_classes.append(new_pred_class)
                inputs.append(new_input_unscaled)

    # Build DataFrame
    data = {col: [] for col in column_names}
    data["Class"] = []
    data["Probability"] = []

    for sample in inputs:
        for i, col in enumerate(column_names):
            data[col].append(sample[i])

    for cls in pred_classes:
        data["Class"].append(cls)

    for prob in pred_probability:
        data["Probability"].append(prob * 100)

    generated_df = pd.DataFrame(data=data)
    return generated_df
```

## Anomaly Score Threshold

The following code decides the anomaly score threshold used to guide generation. The first threshold represents moderately anomalous data whereas the second threshold represents extremely anomalous data.

```

In [313... threshold_1 = np.percentile(scores,5)
threshold_2 = -0.15

In [314... exponents = np.linspace(-12, -2 , 600)
positive_epsilon = 10 ** exponents
negative_epsilon = -positive_epsilon
epsilon = np.sort(np.concatenate([negative_epsilon, positive_epsilon]))

threshold = threshold_2 # Change threshold as needed
results = []

for epsilon in epsilon:
    generated_df = generate_data_with_epsilon(model, epsilon)
    anomaly_score_results = []

    for cls in generated_df["Class"].unique():
        subset = generated_df[generated_df["Class"] == cls]
        features = subset.drop(columns=["Class", "Probability"]).to_numpy()
        score = detector.decision_function(features)

        anomaly_score_results.append({
            "Class": cls,
            "Score": score.mean()
        })

    scored_df = pd.DataFrame(anomaly_score_results)
    mean_score_by_class = scored_df.set_index("Class")["Score"]
    below_threshold = mean_score_by_class[mean_score_by_class < threshold].to_dict()

    results.append({
        "epsilon": epsilon,
        "mean_score": mean_score_by_class.to_dict(),
    })

In [315... class_names = []
for r in results:
    for cls in r["mean_score"]:
        if cls not in class_names:
            class_names.append(cls)
class_names.sort()

# Plot
plt.figure(figsize=(10, 6))
for cls in class_names:
    eps = []
    overlaps = []
    for r in results:
        eps.append(r["epsilon"])
        overlaps.append(r["mean_score"].get(cls))
    plt.plot(eps, overlaps, marker='o', label=f'Class {cls}')


# Threshold Line
plt.axhline(y=threshold, color='red', linestyle='--', label=f'Threshold ({threshold})')

# Log scale on x-axis

plt.xlabel("Epsilon")
plt.ylabel("Mean Anomaly Score (Isolation Forest)")

plt.title("Mean Anomaly Score (Isolation Forest) vs. Epsilon (Alpha = 1)")
plt.legend()
plt.xscale("symlog")
plt.grid(True, which="both", ls="--", linewidth=0.5)
plt.tight_layout()
plt.show()

```



```
In [316... best_epsilons = {}
best_eps_overall = None
max_abs_eps = 0 # Track largest absolute epsilon

for cls in class_names:
    best_pos_eps = None
    best_neg_eps = None
    best_pos_diff = float('inf')
    best_neg_diff = float('inf')

    for r in results:
        score = r["mean_score"].get(cls)
        eps = r["epsilon"]

        if score is not None and score < threshold:
            diff = threshold - score

            if eps > 0 and diff < best_pos_diff:
                best_pos_diff = diff
                best_pos_eps = eps

            elif eps < 0 and diff < best_neg_diff:
                best_neg_diff = diff
                best_neg_eps = eps

    best_epsilons[cls] = {
        "best_positive": best_pos_eps,
        "best_negative": best_neg_eps
    }

    for eps in [best_pos_eps, best_neg_eps]:
        if eps is not None and abs(eps) > max_abs_eps:
            max_abs_eps = abs(eps)
            best_eps_overall = eps

for cls, eps_dict in best_epsilons.items():
    print(f"Class {cls}:")
    print(f" Best positive epsilon: {eps_dict['best_positive']}")
    print(f" Best negative epsilon: {eps_dict['best_negative']}")

print(f"\nBest overall epsilon across all classes: {best_eps_overall}")
```

Class 0:  
 Best positive epsilon: 2.467860087458031e-07  
 Best negative epsilon: -2.285238607695462e-07

Class 1:  
 Best positive epsilon: 2.769516817469001e-07  
 Best negative epsilon: -1.885641372950549e-07

Best overall epsilon across all classes: 2.769516817469001e-07

## Generated Confidently Classified Anomalies

```

In [317... epsilon = max_abs_eps

pred_classes = []
inputs = []
original_inputs = []
original_classes = []
original_prob = []
pred_probability = []

for test,truth in zip(X_test_scaled,y_test_encoded):
    x_sample = test.reshape(-1, 1)
    pred = model.forward(x_sample)

    class_probs = pred[0:2].ravel()
    max_prob = np.max(class_probs)
    if np.argmax(pred[:2]) == np.argmax(truth):
        if max_prob >= 0.8:
            original_prob.append(max_prob)
            pred_class = np.argmax(class_probs)
            original_inputs.append(ANN_scaler.inverse_transform(x_sample.reshape(1, -1))[0])
            original_classes.append(pred_class)

            vector_neg_epsilon = pred.copy()
            vector_neg_epsilon[2:] = vector_neg_epsilon[2:] - epsilon

            # Reverse pass for when epsilon is subtracted to latent variables
            new_input_neg = model.reverse(input=vector_neg_epsilon)
            new_pred = model.forward(new_input_neg)
            max_prob = np.max(new_pred[0:2])
            pred_probability.append(max_prob)
            new_pred_class_neg = np.argmax(new_pred[0:2])

            vector_pos_epsilon = pred.copy()
            vector_pos_epsilon[2:] = vector_pos_epsilon[2:] + epsilon

            # Reverse pass for when epsilon is added to latent variables
            new_input_pos = model.reverse(input=vector_pos_epsilon)
            new_pred = model.forward(new_input_pos)
            max_prob_pos = np.max(new_pred[0:2])
            pred_probability.append(max_prob_pos)
            new_pred_class_pos = np.argmax(new_pred[0:2])

            new_input_unscaled_pos = ANN_scaler.inverse_transform(new_input_pos.reshape(1, -1))[0]
            new_input_unscaled_neg = ANN_scaler.inverse_transform(new_input_neg.reshape(1, -1))[0]

            pred_classes.append(new_pred_class_pos)
            inputs.append(new_input_unscaled_pos)

            pred_classes.append(new_pred_class_neg)
            inputs.append(new_input_unscaled_neg)

data = {col: [] for col in column_names}
data["Class"] = []
data["Probability"] = []

for sample in inputs:
    for i, col in enumerate(column_names):
        data[col].append(sample[i])

for cls in pred_classes:
    data["Class"].append(cls)

for prob in pred_probability:
    data["Probability"].append(prob * 100)

generated_df = pd.DataFrame(data=data)

In [318... generated_df = generated_df.drop_duplicates()

In [319... prediction_list = detector.predict(inputs).tolist()

anomalies = prediction_list.count(-1) / len(inputs)
print(f"Anomaly rate: {anomalies:.2%}")

inputs = np.array(inputs)
predictions = np.array(prediction_list)

anomalous_inputs = inputs[predictions == -1]

```

```
generated_anomalies = pd.DataFrame(anomalous_inputs, columns=column_names)
```

Anomaly rate: 96.82%

```
In [320... generated_anomalies=generated_anomalies.drop_duplicates().reset_index(drop=True)
```

```
In [321... prediction_list = detector.predict(X).tolist()

anomalies = prediction_list.count(-1)/len(X)

print(anomalies)
```

0.06288280930992242

```
In [322... generated_anomalies
```

```
Out[322...
      fixed    volatile    citric    residual    chlorides    free sulfur    total sulfur    density    pH    sulphates    alcohol
      acidity    acidity    acid    sugar                   dioxide    dioxide
0    6.638056  -0.254595   2.005973   61.009207   0.313283  -149.417803   586.931828   1.037689   3.159388  -1.704244   41.136716
1    6.138499   0.818355   1.167119   18.833559   0.155920   49.093370   281.641924   0.989733   2.872475   0.744614   12.146138
2   10.558523   0.332413   0.510640  -25.920354   0.146480  -37.915215   -8.224125   0.996807   3.258488  -0.076455   16.101660
3   15.325635  -0.137105   0.625696   11.516162   0.052455  -101.817652   80.939176   0.977256   3.445755   1.205317   5.829430
4    2.527764  -1.581751  -1.426902    0.437439   0.263298  -258.334009   441.853298   1.035607   4.941588  -4.760838   44.067869
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
1058 16.851784   2.002280   4.835400  119.823058   0.365007  289.560921  1098.417528   1.009379   0.857623   2.422075   12.885632
1059  9.067925   0.198275   2.041501  130.569115   0.306971  -39.819218   936.524656   1.063838   2.648050  -1.731709   38.123577
1060  5.242954  -1.683543  -3.220992  -88.917411  -0.154664  -434.786442  -548.940859   0.943230   6.022245   0.054602   4.820626
1061  3.677648  -0.695023  -0.684513  -11.623281  -0.046505  -113.578102   99.666006   1.020792   4.443418  -1.535666   25.587627
1062  5.469491  -0.928761  -1.710991  -52.786363  -0.107396  -245.385770  -253.856741   0.962619   4.831910   0.354565   7.606839
```

1063 rows × 11 columns

## Transferability Of Confidently Clasified Anomalies

### Testing Settings

```
In [323... confidence_threshold = 0.8
```

### Randomn Forest

```
In [324... RF_model = RandomForestClassifier(n_estimators=200, random_state=42)
RF_model.fit(X_train, y_train)

y_pred = RF_model.predict(X_test)

test_acc = accuracy_score(y_test, y_pred)
test_precision = precision_score(y_test, y_pred, average="macro")
test_recall = recall_score(y_test, y_pred, average="macro")
print("Test Accuracy:", test_acc)
print("Test Precision:", test_precision)
print("Test Recall:", test_recall)
```

Test Accuracy: 0.8908163265306123

Test Precision: 0.8663732677590605

Test Recall: 0.7937426297169812

```
In [325... probs_all = RF_model.predict_proba(generated_anomalies.values)

max_probs = np.max(probs_all, axis=1)
pred_classes = np.argmax(probs_all, axis=1)

mask = max_probs >= confidence_threshold
RF_anomalies_list = generated_anomalies.values[mask]
max_prob_rf = pred_classes[mask]
high_confidence_count = np.sum(mask)
```



```
In [326... robustness = high_confidence_count/len(generated_anomalies.values)
print(robustness)
```

```
0.09407337723424271
```

```
In [327... RF_df = pd.DataFrame(RF_anomalies_list, columns = column_names)
```

```
In [328... len(RF_anomalies_list)
```

```
Out[328... 100
```

## Testing On Neural Networks

```
In [329... np.random.seed(42)
torch.manual_seed(42)
torch.cuda.manual_seed_all(42)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False

# Model architecture
class SimpleNet(nn.Module):
    def __init__(self, input_dim):
        super(SimpleNet, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, 64),
            nn.ReLU(),
            nn.Linear(64, 32),
            nn.ReLU(),
            nn.Linear(32, 16),
            nn.ReLU(),
            nn.Linear(16, 8),
            nn.ReLU(),
            nn.Linear(8, 2)
        )

    def forward(self, x):
        return self.net(x)

X_train_tensor = torch.tensor(X_train_scaled, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.long)
X_test_tensor = torch.tensor(X_test_scaled, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test, dtype=torch.long)

# Initialize model, loss, optimizer
test_model = SimpleNet(input_dim=X_train_tensor.shape[1])
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(test_model.parameters(), lr=0.001)

# Training Loop
num_epochs = 440
for epoch in range(num_epochs):
    test_model.train()
    optimizer.zero_grad()
    outputs = test_model(X_train_tensor)
    loss = criterion(outputs, y_train_tensor)
    loss.backward()
    optimizer.step()

    if epoch % 10 == 0 or epoch == num_epochs - 1:
        test_model.eval()
        with torch.no_grad():
            test_outputs = test_model(X_test_tensor)
            preds = torch.argmax(test_outputs, dim=1)
            acc = accuracy_score(y_test_tensor, preds)
            print(f"Epoch {epoch:3d} | Loss: {loss.item():.4f} | Test Accuracy: {acc * 100:.2f}%")
```

Epoch	Loss	Test Accuracy
0	0.6743	78.37%
10	0.6495	78.37%
20	0.6229	78.37%
30	0.5871	78.37%
40	0.5369	78.37%
50	0.4868	78.37%
60	0.4642	78.37%
70	0.4469	78.37%
80	0.4328	78.37%
90	0.4198	78.37%
100	0.4093	78.37%
110	0.4013	78.98%
120	0.3944	81.12%
130	0.3881	82.35%
140	0.3824	81.12%
150	0.3769	80.82%
160	0.3718	80.82%
170	0.3668	80.31%
180	0.3615	80.82%
190	0.3560	81.12%
200	0.3502	80.61%
210	0.3440	80.51%
220	0.3369	80.82%
230	0.3293	81.33%
240	0.3213	82.04%
250	0.3126	82.04%
260	0.3029	82.65%
270	0.2918	82.86%
280	0.2804	83.37%
290	0.2694	82.76%
300	0.2586	83.16%
310	0.2473	83.27%
320	0.2370	83.98%
330	0.2280	83.67%
340	0.2193	83.67%
350	0.2121	84.18%
360	0.2049	84.49%
370	0.1982	84.08%
380	0.1916	84.39%
390	0.1856	84.08%
400	0.1800	84.80%
410	0.1743	84.90%
420	0.1691	84.29%
430	0.1651	85.20%
439	0.1630	85.41%

In [330...

```
test_model.eval()
correct = 0
y_preds = []
y_true = []

for i in range(X_test_scaled.shape[0]):
    x_sample = torch.tensor(X_test_scaled[i].reshape(1, -1), dtype=torch.float32)
    y_sample = y_test_encoded[i].reshape(-1) # Assuming one-hot encoded

    with torch.no_grad():
        logits = test_model(x_sample)
        probs = torch.softmax(logits, dim=1).numpy().flatten()
        pred_class = np.argmax(probs)
        true_class = np.argmax(y_sample)

    y_preds.append(pred_class)
    y_true.append(true_class)

    if pred_class == true_class:
        correct += 1

x_sample = torch.tensor(X_test_scaled[0].reshape(1, -1), dtype=torch.float32)
with torch.no_grad():
    pred = test_model(x_sample)

test_accuracy = correct / len(X_test_scaled)
precision = precision_score(y_true, y_preds, average='macro', zero_division=0)
recall = recall_score(y_true, y_preds, average='macro', zero_division=0)
f1 = f1_score(y_true, y_preds, average='macro', zero_division=0)

print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
print(f"Precision: {precision * 100:.2f}%")
print(f"Recall: {recall * 100:.2f}%")
print(f"F1 Score: {f1 * 100:.2f}%")
```

Test Accuracy: 85.41%  
 Precision: 78.71%  
 Recall: 77.20%  
 F1 Score: 77.91%

```
In [331... MLP_anomalies_list = []
high_confidence_count = 0
test = []
scaled_anomalous_inputs = ANN_scaler.transform(generated_anomalies.values)

test_model.eval()

for x in scaled_anomalous_inputs:
    x_tensor = torch.tensor(x.reshape(1, -1), dtype=torch.float32) # shape: [1, input_dim]
    with torch.no_grad():
        logits = test_model(x_tensor) # shape: [1, 2]
        probs = torch.softmax(logits, dim=1).numpy().flatten() # convert to numpy array
        max_prob = np.max(probs)

        if max_prob >= confidence_threshold:
            high_confidence_count += 1
            MLP_anomalies_list.append(x)
            if np.argmax(probs) == 1:
                test.append(x)

MLP_anomalies_list = ANN_scaler.inverse_transform(MLP_anomalies_list)

robustness = high_confidence_count / len(generated_anomalies.values)
print(f"Robustness: {robustness}")
```

Robustness: 0.9877704609595485

```
In [332... len(MLP_anomalies_list)
```

Out[332... 1050

```
In [333... MLP_df = pd.DataFrame(MLP_anomalies_list, columns = column_names)
```

## KNN Classifier

```
In [334... from sklearn.neighbors import KNeighborsClassifier

n = 5
neigh = KNeighborsClassifier(n_neighbors=n)
neigh.fit(X_train_scaled, y_train)

y_pred = neigh.predict(X_test_scaled)

test_acc = accuracy_score(y_test, y_pred)
test_precision = precision_score(y_test, y_pred, average="macro")
test_recall = recall_score(y_test, y_pred, average="macro")
print("Test Accuracy:", test_acc)
print("Test Precision:", test_precision)
print("Test Recall:", test_recall)
```

Test Accuracy: 0.8418367346938775  
 Test Precision: 0.7709322843652139  
 Test Recall: 0.7385883451257862

```
In [335... KNN_anomalies_list = []
high_confidence_count = 0

scaled_anomalous_inputs = ANN_scaler.transform(generated_anomalies.values)

for x in scaled_anomalous_inputs:
    probs = neigh.predict_proba(x.reshape(1, -1))
    max_prob = np.max(probs)

    if max_prob >= confidence_threshold:
        high_confidence_count += 1
        KNN_anomalies_list.append(x)

robustness = high_confidence_count / len(scaled_anomalous_inputs)
KNN_anomalies_list = ANN_scaler.inverse_transform(KNN_anomalies_list)
print(robustness)
```

0.8447789275634995

```
In [336... len(KNN_anomalies_list)
```

Out[336... 898

```
In [337... knn_df = pd.DataFrame(KNN_anomalies_list, columns = column_names)
```

## Checking Shared Vulnerabilities

```
In [338... feature_names = column_names.tolist()
```

```
In [339... # Drop duplicates in each DataFrame
df_rf = RF_df.drop_duplicates()
df_knn = knn_df.drop_duplicates()
df_nn = MLP_df.drop_duplicates()

# Merge on all columns to find common rows
common_rows = df_rf.merge(df_knn, how='inner').merge(df_nn, how='inner')

print("Number of common rows:", len(common_rows))
```

Number of common rows: 89

```
In [340... # Convert to NumPy array if needed
inputs = common_rows.values

# Predict with both models
rf_preds = RF_model.predict(inputs)
knn_preds = neigh.predict(inputs)

# Find indices where class == 1
rf_class1_indices = np.where(rf_preds == 1)[0]
knn_class1_indices = np.where(knn_preds == 1)[0]

# Extract corresponding samples
rf_class1_samples = inputs[rf_class1_indices]
knn_class1_samples = inputs[knn_class1_indices]
```

```
In [341... common_rows
```

```
Out[341... 
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	6.619644	0.725416	1.169992	17.344795	0.133857	79.184379	271.801423	0.995876	2.889040	0.734420	9.590628
1	7.131828	0.488843	0.638656	17.297463	0.075126	82.282676	161.853428	0.994377	2.921193	0.709142	9.564306
2	6.438273	0.843791	0.941368	30.852560	0.143286	70.454402	320.134621	0.995497	2.918892	0.616214	10.750064
3	7.770855	0.478410	0.632777	23.874041	0.067491	89.246862	190.458986	0.996373	2.904713	0.743790	8.957409
4	6.955438	0.386481	1.001164	12.443815	0.071601	66.952040	247.900174	0.995215	3.042596	0.581948	9.182088
...	...	...	...	...	...	...	...	...	...	...	...
84	6.112677	0.334243	0.286982	8.328684	0.055494	88.856758	50.597346	0.996917	2.948322	0.611817	7.639219
85	10.797688	0.983184	2.376787	57.817333	0.155730	147.198132	481.909374	1.004500	1.972519	1.271915	9.915220
86	8.328823	0.595884	0.839332	33.656592	0.095419	95.008638	263.783051	0.998548	2.790368	0.723396	9.256311
87	8.349452	0.439609	0.551893	18.302610	0.065321	52.645978	118.127321	0.993491	3.058263	0.628026	9.043823
88	7.816388	0.422726	0.554827	16.202754	0.058204	54.011559	109.120821	0.994899	3.075205	0.765567	7.694477

89 rows × 11 columns

```
In [342... # Filter rows where any feature column has a negative value
negative_rows = common_rows[feature_names][common_rows[feature_names] < 0].any(axis=1)]
```

```
In [343... negative_rows
```

```
Out[343... 
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
19	5.982089	0.199720	0.008589	-3.455079	0.007816	37.162714	99.332774	0.994164	3.467651	0.260337	8.736993
47	4.839145	0.146253	-0.304224	4.579052	0.042293	66.979595	25.338325	0.996992	3.256260	0.522848	5.238934
58	9.446042	0.240549	0.122134	-6.646533	0.092361	6.782285	113.525672	0.998674	3.301367	0.043708	10.669548

```
In [344... negative_rows.reset_index(drop=True)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	5.982089	0.199720	0.008589	-3.455079	0.007816	37.162714	99.332774	0.994164	3.467651	0.260337	8.736993
1	4.839145	0.146253	-0.304224	4.579052	0.042293	66.979595	25.338325	0.996992	3.256260	0.522848	5.238934
2	9.446042	0.240549	0.122134	-6.646533	0.092361	6.782285	113.525672	0.998674	3.301367	0.043708	10.669548

In [345... common\_rows[feature\_names].iloc[14]

Out[345... fixed acidity 7.400440  
volatile acidity 0.575708  
citric acid 1.144452  
residual sugar 23.253191  
chlorides 0.101103  
free sulfur dioxide 92.268847  
total sulfur dioxide 336.973133  
density 0.996635  
pH 2.789534  
sulphates 0.857810  
alcohol 10.159199  
Name: 14, dtype: float64

## Example of Anomalous Sample

In [346... sample\_row = negative\_rows.iloc[1]  
print("Input row (original scale):")  
print(sample\_row)

Input row (original scale):  
fixed acidity 4.839145  
volatile acidity 0.146253  
citric acid -0.304224  
residual sugar 4.579052  
chlorides 0.042293  
free sulfur dioxide 66.979595  
total sulfur dioxide 25.338325  
density 0.996992  
pH 3.256260  
sulphates 0.522848  
alcohol 5.238934  
Name: 47, dtype: float64

In [347... sample\_row = common\_rows.iloc[14] # You can change the index if needed  
  
scaled\_sample = ANN\_scaler.transform(sample\_row.values.reshape(1, -1))  
  
x\_tensor = torch.tensor(scaled\_sample, dtype=torch.float32)  
test\_model.eval()  
  
with torch.no\_grad():  
logits = test\_model(x\_tensor)  
probs = torch.softmax(logits, dim=1).numpy().flatten()  
max\_prob = np.max(probs)  
predicted\_class = np.argmax(probs)  
print("Neural Network Model")  
print("Probability of Low Quality:", probs[0]\*100,"%")  
print("Probability of Hih Quality Quality:", probs[1]\*100,"%")

Neural Network Model  
Probability of Low Quality: 100.0 %  
Probability of Hih Quality Quality: 1.2265634e-20 %

In [348... probs = neigh.predict\_proba(sample\_row.values.reshape(1,-1))  
print("K-NN Model (K = 5)")  
print("Probability of Low Quality:", probs[0][0]\*100,"%")  
print("Probability of Hih Quality Quality:", probs[0][1]\*100,"%")

K-NN Model (K = 5)  
Probability of Low Quality: 100.0 %  
Probability of Hih Quality Quality: 0.0 %

In [349... probs = RF\_model.predict\_proba(sample\_row.values.reshape(1,-1))  
print("Random Forest Model")  
print("Probability of Low Quality:", probs[0][0]\*100,"%")  
print("Probability of Hih Quality Quality:", probs[0][1]\*100,"%")

Random Forest Model  
Probability of Low Quality: 80.5 %  
Probability of Hih Quality Quality: 19.5 %

Probability of Low Quality: 80.5 %  
Probability of Hih Quality Quality: 19.5 %